

```

import pandas as pd
import numpy as np
from enum import Enum
import matplotlib.pyplot as plt

class Constant_DP(Enum):
    ACTION = [i for i in range(21)]
    T = 10
    D = [i for i in range(11)]
    INITIAL_STATES = [-10, -9, -8, -7, -6, -5, -4, -3, -2, -1, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

```

```

def cost_function(s_t, a_t, D_t, o_t, h_t, b_t):
    s_t_plus1 = s_t + a_t - D_t
    return o_t*a_t + max(h_t*(s_t_plus1), -b_t*(s_t_plus1))

```

```

def plot_function(y, ylabel, title):
    # x = [i for i in range(len(y))]
    plt.plot(Constant_DP.INITIAL_STATES.value, y)
    plt.xlabel("Initial States (s_0)")
    plt.ylabel(f"{ylabel}")
    plt.title(f"{title}")
    plt.show()

```

```

print(Constant_DP.INITIAL_STATES.value)

[-10, -9, -8, -7, -6, -5, -4, -3, -2, -1, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

```

```

def compute_possible_states():
    max_demand = np.max(Constant_DP.D.value) # this is min s_t+1
    max_action = np.max(Constant_DP.ACTION.value) # this is max s_t+1
    horizon = 10 # number of time steps
    min_initial_state = -10
    max_initial_state = 10

    return [value for value in range(min_initial_state - horizon*max_demand, horizon*max_action + max_initial_state + 1)]

```

```

def dp_implementation(order, holding, backlog):
    """
    """
    all_possible_states = compute_possible_states()
    value_cost_matrix = [[0 for _ in range(len(all_possible_states))] for t in range(11)] # rows corresponds to time step and columns
    value_cost_matrix[-1] = np.array([max(holding*state, -backlog*state) for state in all_possible_states])
    uniform_probability = 1/len(Constant_DP.D.value)
    policy_matrix = [[0 for _ in range(len(all_possible_states))] for t in range(11)] # final step you take no action, so policy of

    state_to_index_map = {state:idx for idx, state in enumerate(all_possible_states)}

    # backwards iteration
    for time_step in range(10, -1, -1):

        if time_step == 10:
            continue

        for idx, state in enumerate(all_possible_states):
            expected_min_cost = float('inf')
            optimal_action = 0

            for action in Constant_DP.ACTION.value:
                expectation = 0

                for demand in Constant_DP.D.value:

                    value_row = value_cost_matrix[time_step+1]
                    next_state = state + action - demand

                    # makes sure the next state is in range of all possible valid states
                    if next_state < all_possible_states[0]:
                        next_state_idx = 0

                    elif next_state > all_possible_states[-1]:
                        next_state_idx = len(all_possible_states) - 1

```

```

else:
    next_state_idx = state_to_index_map[next_state]
    next_value_cost = value_cost_matrix[time_step+1][next_state_idx]

    # finding the expected cost with the probability of getting the demand
    expectation += uniform_probability*(cost_function(state, action, demand, order, holding, backlog) + next_value_cost)

if expectation < expected_min_cost:
    optimal_action = action
    expected_min_cost = expectation

policy_matrix[time_step][idx] = optimal_action # keeps track of the optimal action
value_cost_matrix[time_step][idx] = expected_min_cost # keeps track of optimal cost

initial_state_indices = [state_to_index_map[initial_state] for initial_state in Constant_DP.INITIAL_STATES.value]
value_cost_initial = [value_cost_matrix[0][initial_state_idx] for initial_state_idx in initial_state_indices]
policy_initial = [policy_matrix[0][initial_state_idx] for initial_state_idx in initial_state_indices]

plot_function(value_cost_initial, "V*0(s_0)", "Plot of V*0(s_0)")
plot_function(policy_initial, "a*0(s_0)", "Plot of a*0(s_0)")

# s*0
optimal_s0_idx = np.argmin(value_cost_initial)
optimal_s0 = Constant_DP.INITIAL_STATES.value[optimal_s0_idx]

# base stock
base_stock_idx = np.argmin(policy_initial)
base_stock = Constant_DP.INITIAL_STATES.value[base_stock_idx]

print("Base stock optimal s_t:", base_stock)
print("The value s0 that minimizes V*0 is: ", optimal_s0)

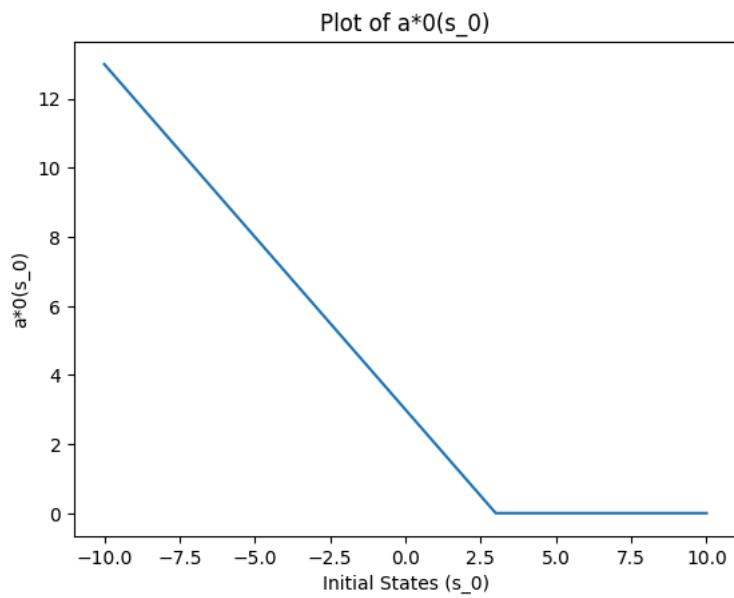
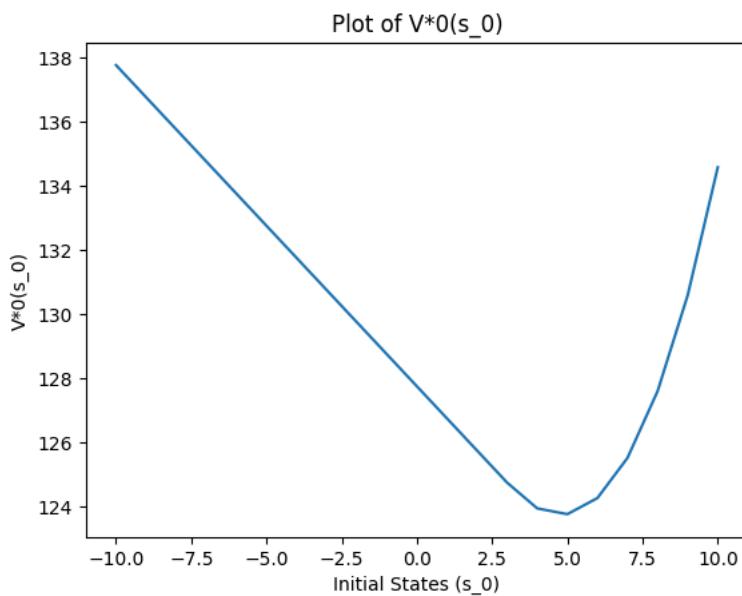
return value_cost_initial, policy_initial

```

```

o_t = 1
h_t = 4
b_t = 2
optimal_value_cost, _ = dp_implementation(o_t, h_t, b_t)

```



Base stock optimal s_t : 3
The value s_0 that minimizes V^* is: 5

```
def find_naive_action(state):
    s_bar = 5
    action = max(0, s_bar-state)

    # returns only valid actions between 0 and 20
    return min(action, 20)
```

```
def dp_naive_implementation(order, holding, backlog):
    """
    """
    all_possible_states = compute_possible_states()
    value_cost_matrix = [[0 for _ in range(len(all_possible_states))]] for t in range(11)] # rows corresponds to time step and columns
    value_cost_matrix[-1] = np.array([max(holding*state, -backlog*state) for state in all_possible_states])
    uniform_probability = 1/len(Constant_DP.D.value)

    state_to_index_map = {state:idx for idx, state in enumerate(all_possible_states)}

    # backwards iteration
    for time_step in range(10, -1, -1):

        if time_step == 10:
            continue

        for idx, state in enumerate(all_possible_states):
```

```

action = find_naive_action(state)
naive_expectation = 0

for demand in Constant_DP.D.value:

    value_row = value_cost_matrix[time_step+1]
    next_state = state + action - demand

    # makes sure the next state is in range of all possible valid states
    if next_state < all_possible_states[0]:
        next_state_idx = 0

    elif next_state > all_possible_states[-1]:
        next_state_idx = len(all_possible_states) - 1

    else:
        # next_state_idx = np.argwhere(all_possible_states == next_state)
        next_state_idx = state_to_index_map[next_state]
        next_value_cost = value_cost_matrix[time_step+1][next_state_idx]

    # finding the expected cost with uniform probability of getting a specific demand
    naive_expectation += uniform_probability*(cost_function(state, action, demand, order, holding, backlog) + next_value_cost)

    value_cost_matrix[time_step][idx] = naive_expectation # keeps track of optimal cost

initial_state_indices = [state_to_index_map[initial_state] for initial_state in Constant_DP.INITIAL_STATES.value]
naive_value_cost_initial = [value_cost_matrix[0][initial_state_idx] for initial_state_idx in initial_state_indices]

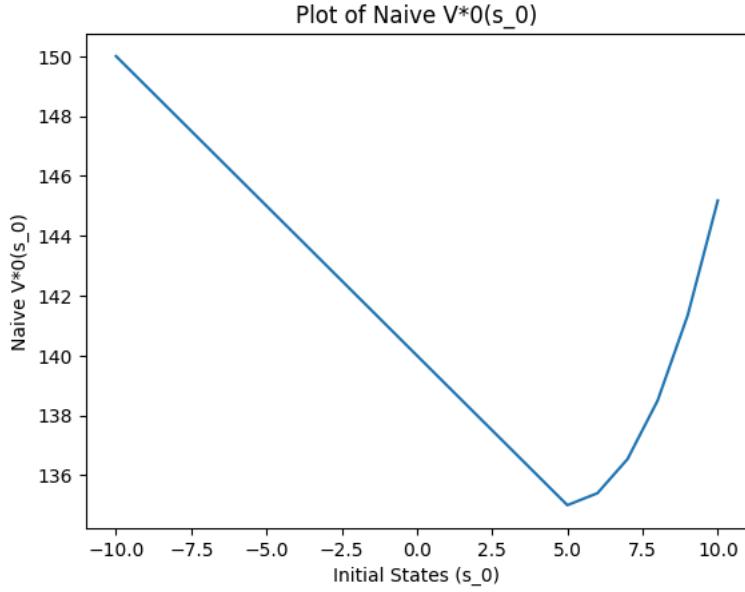
plot_function(naive_value_cost_initial, "Naive V*0(s_0)", "Plot of Naive V*0(s_0)")

# s*0
optimal_s0_idx = np.argmin(naive_value_cost_initial)
optimal_s0 = Constant_DP.INITIAL_STATES.value[optimal_s0_idx]
print("The value s0 that minimizes V*0 is: ", optimal_s0)

return naive_value_cost_initial, []

```

```
naive_value_cost, _ = dp_naiveImplementation(o_t, h_t, b_t)
```



The value s_0 that minimizes V^0 is: 5

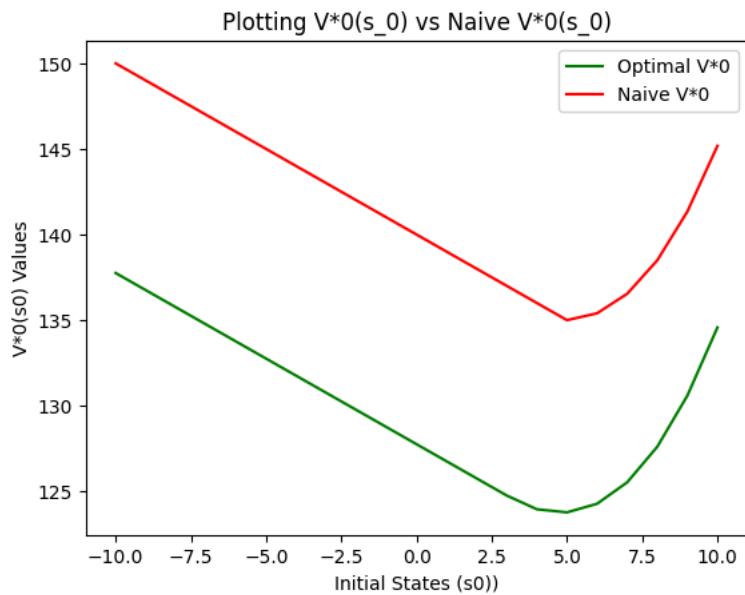
```

value_cost_comparison_array = [optimal_value_cost, naive_value_cost]
colors = ["green", "red"]
labels = ["Optimal  $V^0$ ", "Naive  $V^0$ "]

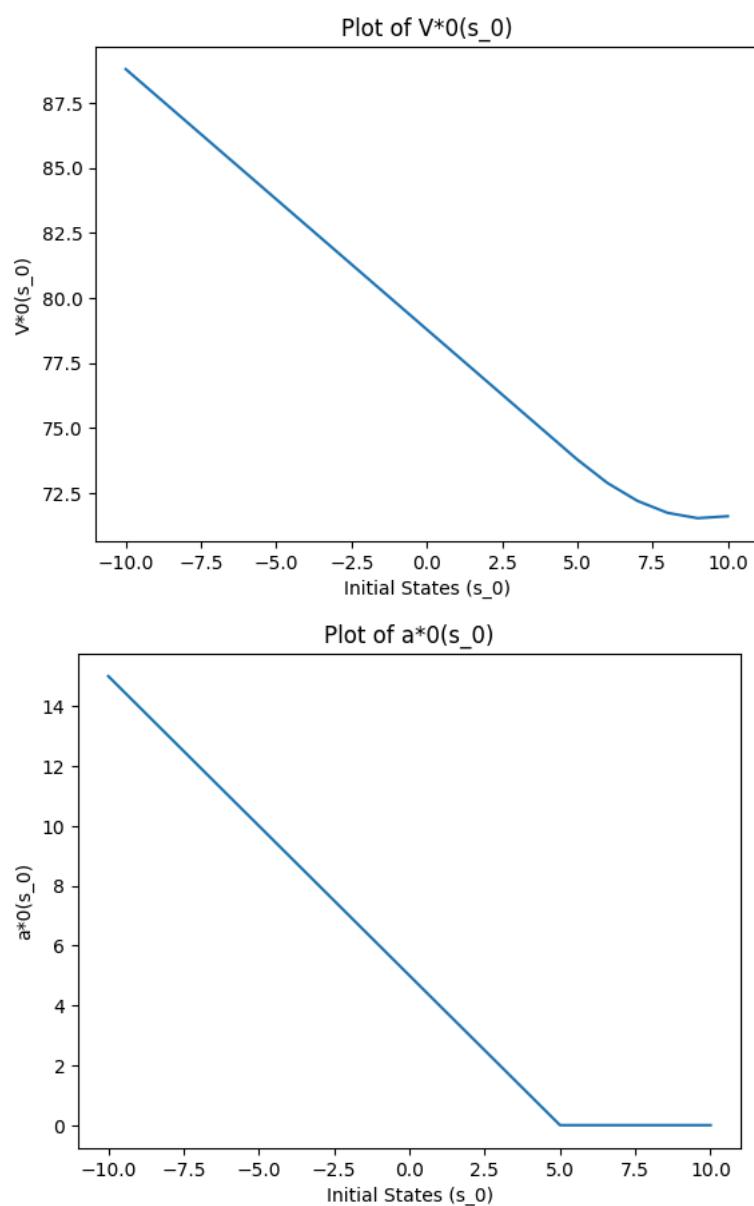
for idx, val_cost in enumerate(value_cost_comparison_array):
    color = colors[idx]
    plt.plot(Constant_DP.INITIAL_STATES.value, val_cost, color, label=labels[idx])
    plt.title("Plotting  $V^0(s_0)$  vs Naive  $V^0(s_0)$ ")
    plt.xlabel("Initial States ( $s_0$ )")
    plt.ylabel("V*0( $s_0$ ) Values")

```

```
plt.legend()  
plt.show()
```

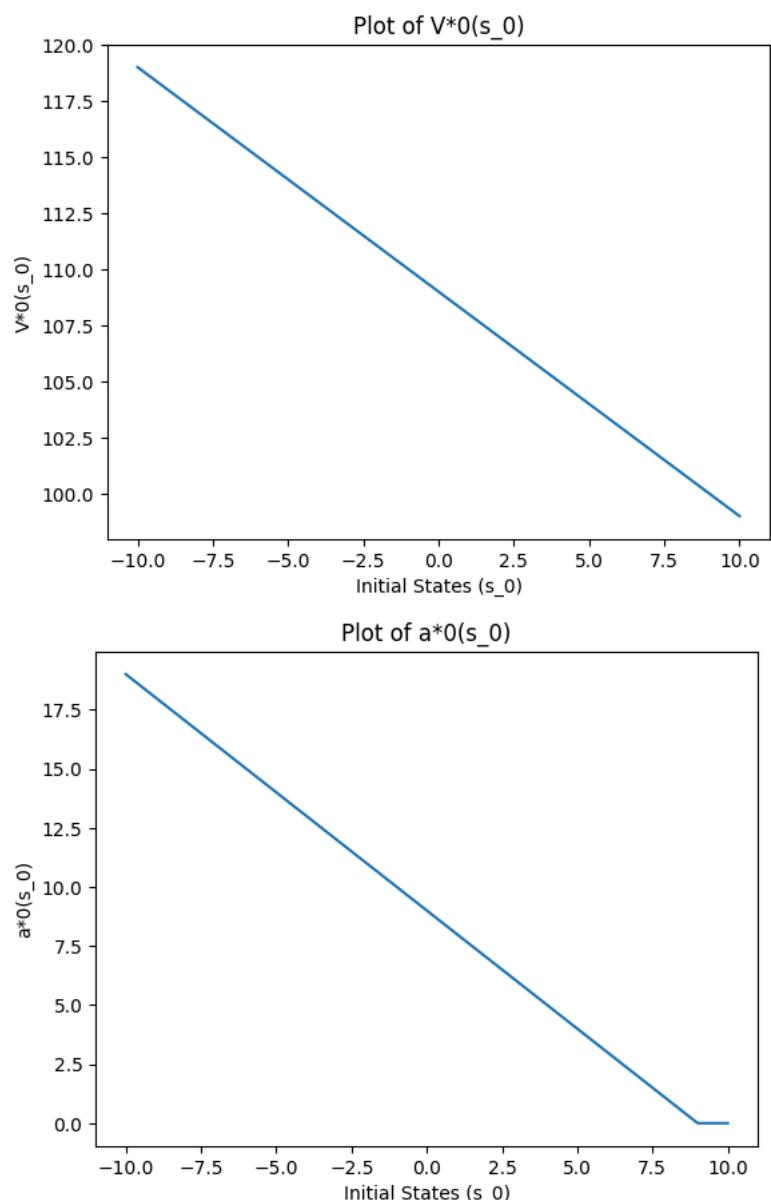


```
# baseline  
o_base = 1  
h_base = 1  
b_base = 1  
base_value_cost, base_policy = dp_implementation(o_base, h_base, b_base)
```



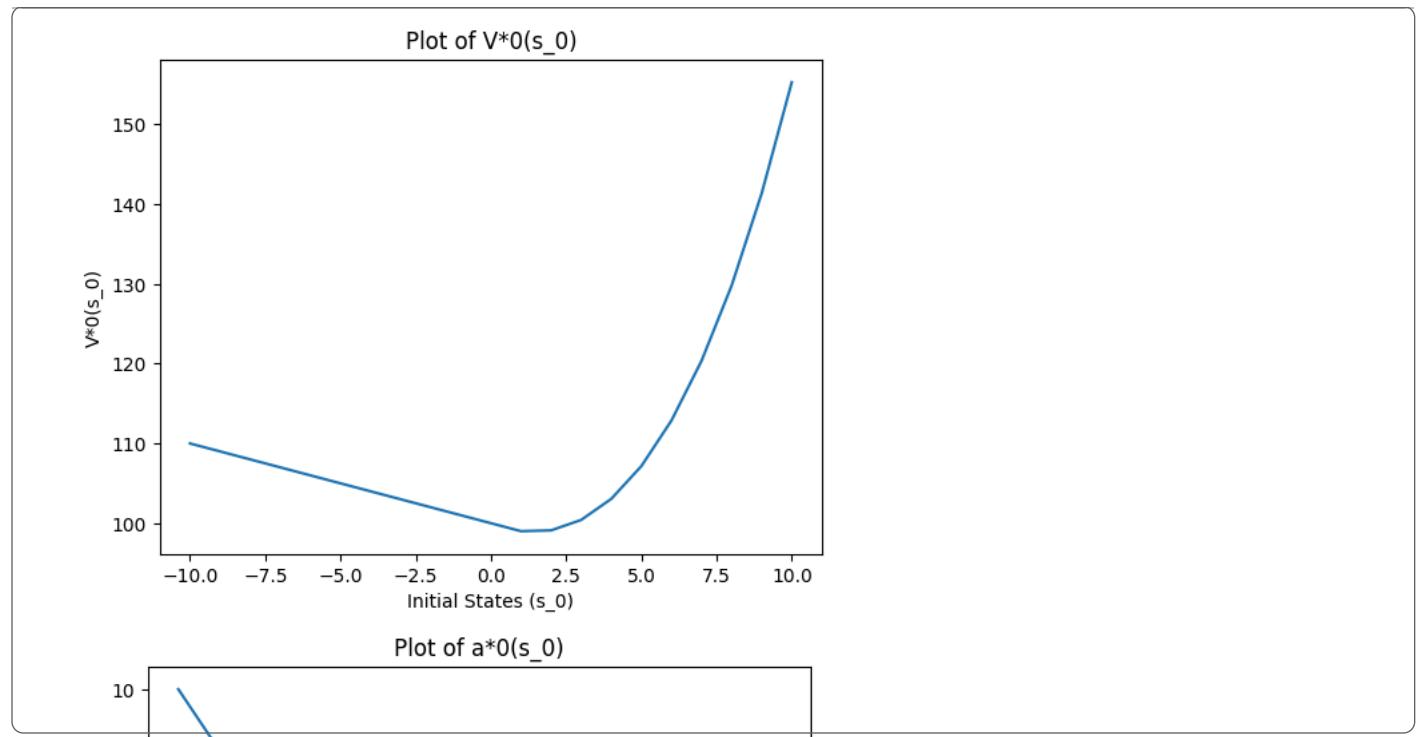
Base stock optimal s_t : 5
The value s_0 that minimizes V^* is: 9

```
# high backlog
b_high = 10
high_b_value_cost, high_b_policy = dpImplementation(o_base, h_base, b_high)
```



Base stock optimal s_t : 9
The value s_0 that minimizes V^* is: 10

```
# high holding
h_high = 10
high_h_value_cost, high_h_policy = dp_implementation(o_base, h_high, b_base)
```



```
# high order
o_high = 11
high_o_value_cost, high_o_policy = dp_implementation(o_high, h_base, b_base)
```

